WhyLab: A Causal Safety Monitoring Framework for Stable Self-Improving Agents

Anonymous Author(s)
Affiliation
Address
email

Abstract

Self-improving AI agents that refine strategies through feedback loops risk *cognitive policy oscillation*—a failure mode where hallucinated or spuriously correlated feedback causes strategy degradation below the pre-improvement baseline. Existing causal inference tools (DoWhy, EconML) provide rigorous offline estimation but do not directly address the need for online safety monitoring and drift-aware update control in noisy agent environments.

We present WhyLab, a causal safety monitoring framework for stable agent self-improvement, comprising three main components: **(C1)** an information-theoretic drift index using entropy-inverse weighting that detects distributional shifts earlier than standard monitoring baselines; **(C2)** a sensitivity filter combining approximate E-values and partial R^2 bounds that reduces false attribution of causal effects under hidden confounders; and **(C3)** a Lyapunov-bounded adaptive damping controller with a sufficient condition for energy non-increase ($\mathbb{E}[\Delta V] \leq 0$).

Controlled experiments on synthetic agent environments and a ReAct-style LLM coding agent (HumanEval) show that the framework improves drift detection reliability, filters fragile successes (reducing oscillation by 16% with minimal pass-rate cost), and stabilizes proxy-observable updates under noisy feedback with hallucination rates up to 50%. Furthermore, we establish a pilot-holdout calibration method to adapt observational sensitivity thresholds to low-sample online agent loops, characterizing the precision-safety Pareto tradeoff.

1 Introduction

Autonomous AI agents that iteratively refine their own strategies through feedback loops—via RAG knowledge updates, prompt mutation, or reinforcement from user signals—represent a powerful paradigm for continuous improvement Shinn et al. [2023], Yao et al. [2023]. However, this self-improvement introduces a critical failure mode: when hallucinated or spuriously correlated feedback accumulates, the agent's policy oscillates between contradictory strategies, a phenomenon we term *cognitive policy oscillation*. Left unchecked, this oscillation can silently degrade performance below the pre-improvement baseline.

The challenge is threefold. **First**, detecting when the agent's operational environment has shifted requires distinguishing genuine distributional drift from normal variance—a signal-to-noise problem that uniform monitoring approaches handle poorly (§3.1). **Second**, not all observed improvements are causally robust; some succeed due to spurious confounders that will not persist, leading to *fragile successes* that inflate confidence (§3.2). **Third**, even with accurate drift detection and sensitivity analysis, the agent's strategy update step size must be bounded to prevent oscillation—but overly

Submitted to 39th Conference on Neural Information Processing Systems (NeurIPS 2025). Do not distribute.

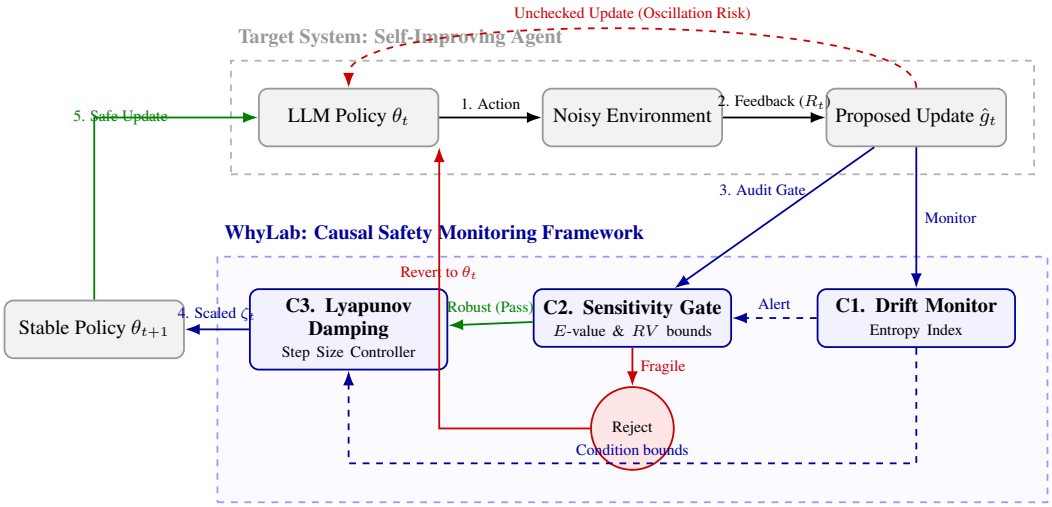

Figure 1: WhyLab causal safety monitoring framework. Instead of uncritically accepting potentially hallucinated updates ($\hat{g}_t$) that cause cognitive oscillation, proposed updates are routed through a safety monitoring layer. **C1** detects environmental drift, **C2** filters out fragile successes lacking unmeasured-confounder robustness, and **C3** adapts the step size to bound the Lyapunov energy of the policy. Unsafe updates are rejected, preserving the status quo.

conservative bounds cause learning deadlock, particularly under heavy-tailed noise typical of LLM evaluation environments (§3.3).

Existing causal inference tools address parts of this problem in isolation. DoWhy Sharma and Kiciman [2022] provides a principled model–identify–estimate–refute workflow, and EconML Battocchi et al. [2019] offers heterogeneous treatment effect estimators including Double Machine Learning. However, these libraries are designed for offline observational data and do not provide an online safety monitoring framework capable of real-time drift detection and bounded strategy updates for continuously operating agents.

We present WhyLab, a causal safety monitoring framework that provides the missing runtime safeguards for stable agent self-improvement. Our framework makes the following contributions:

- **C1. Information-theoretic drift detection** (§3.1). An entropy-weighted drift index that automatically adapts signal importance based on information content, providing an early warning system for environmental shifts.
- **C2. Sensitivity-aware effect filtering** (§3.2). Integration of E-values and partial R^2 bounds to quantify robustness against unmeasured confounders, acting as a safety gate to filter fragile successes before they trigger harmful updates.
- **C3. Lyapunov-bounded adaptive damping** (§3.3). A mathematically grounded controller with a sufficient condition for energy non-increase, translating drift severity (C1) and sensitivity confidence (C2) into safe, bounded policy steps.

We validate WhyLab on both controlled synthetic environments and an actual self-improving LLM coding agent based on Reflexion Shinn et al. [2023]. Across both regimes, the safety monitoring framework successfully prevents cognitive oscillation with minimal cost to overall performance, confirming its practical utility for safe agent deployment.

2 Problem Formulation

Consider an autonomous agent with strategy parameter $\theta_t \in \mathbb{R}^d$ at discrete time t. The agent observes a reward signal R_t and updates its strategy via:

$$\theta_{t+1} = \theta_t - \zeta_t \hat{g}_t \tag{1}$$

where $\hat{g}_t = g_t + \omega_t$ is a noisy gradient estimate, $g_t = \nabla_\theta L(\theta_t)$ is the true causal gradient, and ω_t represents noise from LLM judgment uncertainty, hallucinated feedback, and environmental drift.

Failure mode. When ω_t is large relative to g_t (e.g., due to hallucination), the update (1) may move θ_t *away* from the optimum θ^*. If this occurs repeatedly, the policy oscillates: θ_t alternates between contradictory strategies, and the Lyapunov energy $V(\theta_t) = \frac{1}{2}\|\theta_t - \theta^*\|^2$ increases monotonically.

Desiderata. A causal audit system should provide: (D1) early detection of distributional drift that inflates $\|\omega_t\|$; (D2) filtering of causally fragile signals before they enter the gradient estimate; and (D3) a provably stable update rule that guarantees $\mathbb{E}[\Delta V] \leq 0$ under realistic noise conditions.

3 Method

3.1 C1: Information-Theoretic Drift Index

Definition 1 (Drift Index)**.** *Given K audit signal streams $S_1, \ldots, S_K$ (e.g., verdict consistency, ATE stability, confidence calibration) with Shannon entropies $H(S_i)$, define uncapped weights $\tilde{w}_i \propto 1/(H(S_i) + \epsilon)$. The final weights are obtained by projecting onto the capped simplex:*

$$w = \Pi_{\{w \geq 0,\ \sum_i w_i = 1,\ w_i \leq w_{\max}\}}(\tilde{w}) \tag{2}$$

with $w_{\max} = 0.6$ and $\epsilon = 0.1$. The drift index is:

$$DI = \sum_{i=1}^{K} w_i \cdot d_i \tag{3}$$

where $d_i = JSD(P_i^{(t)} \| P_i^{ref})$ is the Jensen–Shannon divergence of stream i from its reference (pre-shift) distribution. We fix k=10 histogram bins per stream for reproducibility. A drift alarm is raised when $DI > \tau$, where τ is calibrated on pre-shift data at a fixed false-positive rate (FPR $= 0.05$).

Intuition. Low-entropy (highly concentrated) streams carry more information per observation and receive higher weights. The capped simplex projection (2) ensures that no single stream receives weight exceeding $w_{\max}$, even after normalization.

Proposition 1 (Detection advantage under inactive projection)**.** *Assume the projection* (2) *is inactive (i.e., all uncapped weights satisfy $\tilde{w}_i \leq w_{\max}$). Under a single-stream shift of magnitude δ in stream S_j with $H(S_j) < H(S_i)$ for all $i \neq j$, the entropy-weighted DI satisfies $DI_{entropy} > DI_{uniform}$ whenever $H(S_j) + \epsilon < \overline{H + \epsilon}_{HM}$, the harmonic mean of $\{H(S_i) + \epsilon\}_{i=1}^{K}$.*

3.2 C2: Sensitivity-Aware Effect Filtering

Not all positive treatment effects are causally robust. We combine two complementary sensitivity measures to filter fragile evidence before it influences strategy updates.

Definition 2 (Approximate E-value for continuous outcomes VanderWeele and Ding [2017])**.** *For a continuous outcome, we compute an approximate E-value using the standard mean-difference conversion VanderWeele and Ding [2017]: the standardized effect $d = ATE/\sigma_{pooled}$ (in the agent context, ATE is the observed episode effect ΔR_t) is converted to an approximate risk-ratio scale via $RR \approx \exp(0.91 \cdot d)$, yielding:*

$$E = RR + \sqrt{RR \cdot (RR - 1)} \tag{4}$$

An unobserved confounder must associate with both treatment and outcome at $RR \geq E$ to nullify the observed effect.

Robustness Value. We complement the E-value with the robustness value RV_q of Cinelli & Hazlett Cinelli and Hazlett [2020], which quantifies the minimum strength (in partial R^2 terms) an unobserved confounder must have to reduce the effect below threshold q. The formal definition is given in Appendix I.

Filtering rule. A decision is classified as *robust* only if $E > E_{\min}$ (default 2.0) *and* $RV_q > RV_{\min}$ (default 0.1). Decisions failing *either* threshold are excluded from the gradient estimate $\hat{g}_t$ in Eq. (1), reducing the effective noise in ω_t.

3.3 C3: Lyapunov-Bounded Adaptive Damping

We derive a sufficient condition for stable strategy updates.

Assumptions.

(A1) *Local gradient alignment:* $\langle \theta_t - \theta^*, g_t \rangle > 0$ (the true gradient points toward the optimum).

(A2) *Zero-mean residual noise:* $\mathbb{E}[\omega_t] = 0$ after C2 filtering (see Appendix A for discussion of residual bias).

(A3) *Finite second moment:* $\mathbb{E}[\|\omega_t\|^2] < \infty$.

Theorem 1 (Sufficient condition for stable updates)**.** *Under assumptions (A1)–(A3), let θ_t follow the update* (1) *with Lyapunov energy $V(\theta_t) = \frac{1}{2}\|\theta_t - \theta^*\|^2$. Then $\mathbb{E}[\Delta V] \leq 0$ if:*

$$\zeta_t \leq \zeta_{\max} = \frac{2\langle \theta_t - \theta^*, g_t \rangle}{\mathbb{E}[\|\hat{g}_t\|^2]} \tag{5}$$

Proof sketch. Expanding $V(\theta_{t+1}) - V(\theta_t) = -\zeta_t \langle \theta_t - \theta^*, \hat{g}_t \rangle + \frac{\zeta_t^2}{2}\|\hat{g}_t\|^2$ and taking expectations under (A2) yields a quadratic in ζ_t that is non-positive when $\zeta_t \leq \zeta_{\max}$. Full proof in Appendix A.

Connection to C1 and C2. The denominator $\mathbb{E}[\|\hat{g}_t\|^2] = \|g_t\|^2 + \mathbb{E}[\|\omega_t\|^2]$ increases when drift is high (large DI from C1) or evidence is fragile (large sensitivity uncertainty from C2). Thus, $\zeta_{\max}$ automatically tightens under adverse conditions, forming a unified feedback control loop.

Practical controller. Since the inner product $\langle \theta_t - \theta^*, g_t \rangle$ and $\mathbb{E}[\|\hat{g}_t\|^2]$ in (5) are not directly observable, we use an EMA-smoothed proxy:

$$\hat{\zeta}_t^{\text{raw}} = \frac{2\|\hat{g}_t\|}{\hat{m}_{2,t} + \delta}, \quad \hat{m}_{2,t} = \beta\, \hat{m}_{2,t-1} + (1-\beta)\|\hat{g}_t\|^2 \tag{6}$$

$$\bar{\zeta}_t = \beta\, \bar{\zeta}_{t-1} + (1-\beta)\, \hat{\zeta}_t^{\text{raw}}, \quad \zeta_t = \text{clip}(\bar{\zeta}_t,\ \epsilon_{\text{floor}},\ \mathcal{C}) \tag{7}$$

where $\hat{m}_{2,t}$ is the EMA second-moment estimate of $\|\hat{g}_t\|^2$, β=0.9, and δ=10^{-10}. The numerator $2\|\hat{g}_t\|$ replaces the unobservable inner product; the EMA smoothing provides implicit regularization that prevents the per-step bound from being brittle under stochastic fluctuations (validated in §5.3).

Floor-ceiling constraint. When noise ω_t is heavy-tailed (e.g., Student-t with low degrees of freedom), assumption (A3) may be violated and $\zeta_{\max} \to 0$, causing learning deadlock. As a practical safeguard, Eq. (7) enforces:

$$\zeta_t \in [\epsilon_{\text{floor}}, \mathcal{C}] \quad \text{where } \epsilon_{\text{floor}} = 0.01,\ \mathcal{C} = 0.8 \tag{8}$$

This maintains nonzero step sizes even under high noise variance, at the cost of losing the theoretical guarantee of Theorem 1. We analyze this trade-off in Appendix D.

Observable energy proxy. Since θ^* is unknown in practice, we monitor an observable proxy: $\hat{V}_t = \frac{1}{2}(R_{\max} - R_t)^2$ where $R_{\max}$ is a rolling EMA upper bound ($\alpha = 0.1$). Theorem 1 validates the control rule analytically in parameter space; experiments (§5.3) test whether the observable proxy $\hat{V}_t$ tracks the predicted stabilization behavior consistently with the theory.

4 System Overview

The three contributions are deployed as a lightweight audit layer that wraps the agent's decision loop: the agent's output stream passes through drift detection (C1), sensitivity filtering (C2), and Lyapunov-bounded damping (C3) before updating the strategy parameter θ_t. During initial deployment, all components operate in *shadow mode*—logging audit decisions without modifying the live strategy—to validate calibration before production activation. A transactional outbox ensures zero data loss under API quota throttling, and a cost-aware circuit breaker caps evaluation spend. Implementation details are in the supplementary material.

5 Experiments

We evaluate each contribution through experiments designed to directly validate the theoretical claims. Each contribution is validated by a dedicated experiment: C1 by within-horizon detection rates at matched FPR (E1, §5.1), C2 by fragile-rate reduction and fragile-success rejection (E2,

Table 1: Within-horizon detection rate (%) at matched FPR $= 5\%$, across 40 seeds (bootstrap 95% CI in brackets). Higher is better. ADWIN uses a speed-optimized variant (single-split budget-bounded scan) rather than the full adaptive-window algorithm of Bifet and Gavaldà [2007].

Severity	**C1 (ours)**	**B4: Uniform**	**B5: ADWIN**
Mild	2.5 [0, 8]	0.0	0.0
Moderate	**20.0** [8, 33]	12.5 [3, 23]	2.5 [0, 8]
Severe	**40.0** [25, 55]	32.5 [18, 48]	30.0 [15, 45]

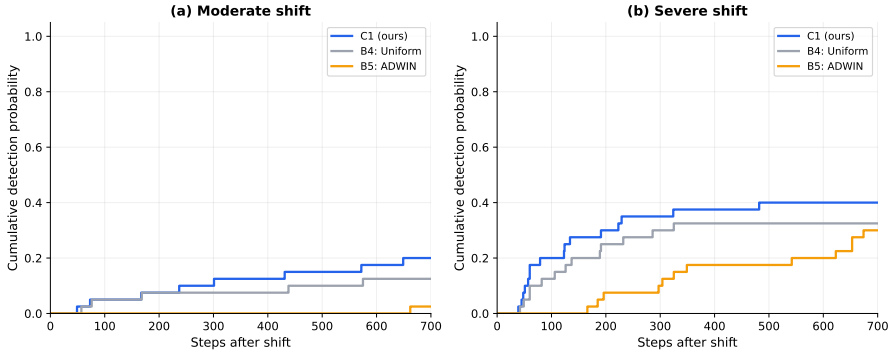

Figure 2: Kaplan–Meier cumulative detection curves for severe shifts (40 seeds). C1 (solid blue) detects shifts earlier and more frequently than uniform (dashed orange) and ADWIN (dotted green). Lower $S(t)$ indicates faster, more reliable detection.

§5.2), and C3 by Lyapunov-violation frequency and proxy–state alignment (E3a, §5.3; heavy-tail analysis in Appendix D).

5.1 E1: Drift Detection Performance

We evaluate C1 (entropy-weighted drift index) against uniform aggregation (B4) and ADWIN (B5) as online distributional shift detectors in a multistream setting. Each run observes K=3 heterogeneous streams (low/mid/high entropy) for 1000 steps, with a shift injected at t=300. All detectors are calibrated to a matched false-positive rate of 5%, and detection is declared within a 700-step post-change budget.

At matched FPR, C1 achieved higher within-horizon detection rates than both baselines across all evaluated severities (Table 1). Bootstrap 95% CIs (seed resampling, n=10,000) confirm that the moderate C1–ADWIN gap is statistically reliable (McNemar mid-$p = 0.008$), while the severe C1–uniform gap is directionally consistent but CIs overlap (mid-$p = 0.125$; McNemar details in supplementary). Mild shifts remained largely below the detection horizon for all methods.

The Kaplan–Meier curves (Figure 2) confirm that C1 detects severe shifts both earlier and more frequently than uniform and ADWIN baselines. Due to conservative thresholding (FPR $= 5\%$), a large fraction of runs remained censored under mild shifts; censoring analysis is provided in Appendix C.2.

5.2 E2: Sensitivity Filtering

We test whether the dual-threshold filter (C2) reliably separates genuine from fragile causal estimates. We generate a synthetic observational study with a known hidden confounder U that inflates the naïve treatment effect. For each of 40 seeds, we estimate ATE via OLS, compute E-values using Def. 2, and compute the robustness value RV_q following Appendix I. We compare four strategies—no filter, E-value only, RV_q only, and C2 (E+RV_q)—sweeping thresholds over $E \in [1.0, 5.0]$ and $RV_q \in [0.01, 0.5]$.

The two axes capture complementary information: E-values measure effect *magnitude* (standardized by outcome variability σ_{pooled}), while RV_q measures statistical *precision* (normalized by regression

Table 2: Filtering performance under two operating regimes from the threshold sweep (40 seeds, mean over synthetic observational studies). For each regime, the operating point minimizing fragile rate is selected per filter. E-values: $d = |ATE|/\sigma_{\text{pooled}}$; RV_q: $f_q = |\hat{\beta}|/\text{SE}$.

Regime	Filter	Frag. rate ↓	Recall ↑	Frag. rej. ↑	Retention
Recall≥0.9	No filter	0.403	1.000	0.000	1.000
	E-only (E≥1.5)	0.335	0.963	0.312	0.813
	RV_q-only (RV_q≥0.50)	0.287	0.939	0.475	0.718
	C2 (E≥1.0, RV_q≥0.50)	**0.287**	0.939	**0.475**	0.718
Recall≥0.5	No filter	0.403	1.000	0.000	1.000
	E-only (E≥2.5)	0.134	0.677	0.873	0.387
	RV_q-only (RV_q≥0.50)	0.287	0.939	0.475	0.718
	C2 (E≥2.5, RV_q≥0.15)	**0.134**	0.677	**0.873**	0.387

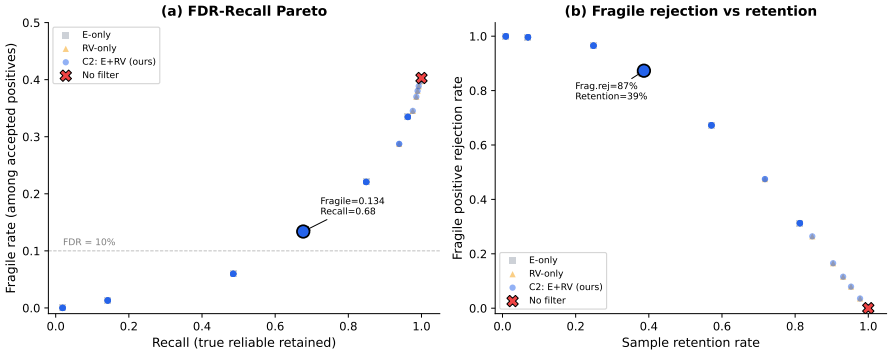

Figure 3: E2 threshold sweep: (a) fragile rate vs. recall Pareto frontier and (b) fragile rejection vs. retention. Each marker represents one (E, RV_q) threshold pair.

SE). Table 2 reports optimal operating points under two recall regimes, applying the same selection rule (minimize fragile rate) to all filters. Under the high-recall regime (recall ≥ 0.9), the RV_q axis dominates: RV_q-only and C2 both reduce fragile rate from 33.5% (E-only) to 28.7%, demonstrating that RV_q captures fragile estimates that E-value alone misses. Under the high-precision regime (recall ≥ 0.5), the E-value axis dominates: E-only and C2 both achieve 13.4% fragile rate and 87.3% fragile rejection. Across both regimes, C2 matches or exceeds the better of the two single-axis filters, confirming that the dual-threshold design never degrades performance while providing a *robustness envelope* that guards against the failure mode of either axis alone.

Refutation tests. To confirm that the robust/fragile separation is not an artifact of threshold sweeping, we conduct three standard refutation tests following the DoWhy workflow Sharma and Kiciman [2022]: (i) adding a random common cause $Z \sim \mathcal{N}(0,1)$ to the estimator, (ii) replacing treatment with a placebo $T_{\text{rand}} \sim \text{Bernoulli}(0.5)$, and (iii) re-estimating on five 80% data subsets. Reliable scenarios pass all three refuters at high rates (data subset 98.6%, random CC 97.4%), while fragile scenarios fail disproportionately (placebo pass rate 28.9–36.2%). The full per-scenario breakdown is in Appendix F, Table 7. These results confirm that C2's causal signals withstand standard refutation perturbations, addressing the identification concern raised for observational causal claims VanderWeele and Ding [2017], Cinelli and Hazlett [2020].

5.3 E3a: Stability Validation

We validate the Lyapunov stability framework (C3) in a stationary quadratic environment with fixed optimum $\theta^*=0$ and $A=\text{diag}(1,2,3,4,5)$, matching assumptions (A1)–(A3). Four controllers are compared: *(B1) Fixed step* $\zeta=0.3$ (tuned to $\zeta < 2/\lambda_{\max}=0.4$); *(B2) Cosine decay* from 0.3 with floor 0.01; *(B3) Gradient clipping* $\zeta_t = \min(0.3,\ c/\|\hat{g}_t\|)$; and the *EMA-smoothed adaptive controller (C3)* from Eq. (7). Noise is zero-mean with hallucination rates $h \in \{0, 0.1, 0.3, 0.5\}$, 40 seeds per condition.

Table 3: Stationary stability metrics with bootstrap 95% CI (seed resampling). Columns 2–4: worst case h=0.5 (40 seeds). Last column: mean across all h values. *Viol. rate*: $P(\Delta V_t > 0)$.

Controller	**Viol. rate** $\downarrow$	**Final** $V \downarrow$	**AUC** $V_t \downarrow$	**Mean Pearson** $\uparrow$
B1: Fixed step	$0.403_{\pm.004}$	5.13	3.78	$0.56_{\pm.01}$
B2: Cosine decay	$0.380_{\pm.004}$	**0.15**	2.24	$0.53_{\pm.02}$
B3: Grad clip	$0.469_{\pm.004}$	0.68	**0.58**	$0.18_{\pm.05}$
C3: Adaptive	$\mathbf{0.328}_{\pm.005}$	3.62	3.25	$\mathbf{0.81}_{\pm.01}$

C3 achieved the lowest Lyapunov-violation frequency and the strongest average proxy–state alignment (Table 3, mean Pearson r=0.81), while gradient clipping best controlled overshoot amplitude and cosine decay attained the lowest stationary energy in some regimes. These results reveal complementary controller profiles rather than uniform dominance: C3 is best positioned as an *online, proxy-observable, noise-aware* stabilizer.

Component ablation. To verify that each component of the practical controller (Eq. 7) is necessary rather than arbitrary, we run a systematic ablation over 6 variants $\times$ 4 hallucination rates $\times$ 20 seeds (Appendix G, Table 8). Removing EMA smoothing collapses violation control to near-random levels (≈ 0.50 across all h), confirming that EMA acts as a *necessary regularizer* for stable bound estimation—not merely a convenience. Drift index (DI from C1) drives the majority of clipping events (removal reduces clip rate from 17.1% to 4.7%), consistent with the intended "tighten under drift" feedback coupling between C1 and C3. Floor and ceiling guards are inactive under moderate noise but prevent extreme final-V values at high hallucination rates, validating the practical safeguard design (§3.3).

An unsmoothed plug-in controller that uses the instantaneous $\zeta_{\max}$ bound diverged (final $V > 1000$), confirming that EMA smoothing acts as an implicit regularizer for the bound estimate (Appendix C.1).

Proxy energy trajectories (Appendix C.1, Figure 4) show that C3-EMA exhibits fewer energy-increasing steps than all baselines at h=0.5.

5.4 Heavy-Tail Robustness

Appendix D analyzes a practical trade-off between raw and EMA-smoothed proxy controllers under Gaussian and heavy-tailed noise. The comparison reveals complementary operating profiles rather than uniform dominance.

Estimator-agnostic operation. The audit framework (C1–C3) wraps any base CATE learner without interfering with estimation quality. We verified compatibility with LinearDML, T-Learner, S-Learner, and X-Learner on the IHDP (1000 replications) and ACIC benchmarks: filter integration did not degrade $\sqrt{\mathrm{PEHE}}$ on any estimator, and method rankings remained invariant (full results in Appendix C). This confirms that C2 operates as a *transparent audit layer*— it screens fragile updates but does not distort the underlying effect estimates.

5.5 E4: Agent Benchmark

To validate that the synthetic-environment results transfer to an actual LLM self-improvement loop, we integrate the WhyLab audit layer into a Reflexion-style agent Shinn et al. [2023] on HumanEval Chen et al. [2021] (164 problems, 5 seeds, 7 ablation configs). The agent iterates *solve* $\rightarrow$ *test* $\rightarrow$ *reflect* $\rightarrow$ *retry* (up to 5 attempts), with each proposed code update gated by C1–C3.

Calibration protocol. E-values were originally proposed as *reporting* tools for observational-study robustness VanderWeele and Ding [2017], not as fixed-threshold automated gates. When applied to an online agent loop with at most 5 attempts and single-problem score windows, the observational-study default thresholds ($E \geq 2.0$, $RV_q \geq 0.1$) reject 85% of proposed updates, causing the audit layer to *degrade* pass-rate below the unaudited baseline. To address this, we conduct a pilot-holdout calibration: a 10-problem pilot set determines environment-adapted thresholds ($E \geq 1.5$, $RV_q \geq 0.05$), corresponding to a shift from "strong" to "moderate" robustness level in the sensitivity taxonomy VanderWeele and Ding [2017]; all reported results are then evaluated on the remaining

Table 4: E4: Reflexion agent on HumanEval (164 problems $\times$ 5 seeds). Pass@1 = final pass rate; Safe = passed with zero regressions; Osc. = oscillation index (pass$\leftrightarrow$fail transitions / max, lower is better); Acc. = update acceptance rate (#accepted / #proposed, N/A where no gate is active). *Default*: observational thresholds ($E \geq 2.0$, $RV_q \geq 0.1$). *Calibrated*: environment-adapted thresholds ($E \geq 1.5$, $RV_q \geq 0.05$) from pilot-holdout calibration.

Ablation	Pass@1	Safe	Osc.	Acc.
none	0.955	0.955	0.051	N/A
C1 only	0.954	0.954	0.051	1.00
C3 only	0.955	0.955	0.051	1.00
Default thresholds ($E \geq 2.0$, $RV_q \geq 0.1$)				
C2 only (default)	0.933	0.933	0.043	0.15
full (default)	0.933	0.933	0.043	0.15
Calibrated thresholds ($E \geq 1.5$, $RV_q \geq 0.05$)				
C2 only (calibrated)	0.955	0.955	0.043	1.00
full (calibrated)	0.955	0.955	0.043	1.00
none (stress)	0.967	0.967	0.053	N/A
full (stress, cal.)	0.967	0.967	0.044	1.00

holdout problems with thresholds frozen. We report *both* operating points (Table 4) to make the precision–safety tradeoff fully transparent.

Results. Without auditing ("none"), the agent achieves 95.5% pass@1. Under the *default* observational thresholds, C2 rejects 85% of proposed updates (Acc. $= 0.15$), yielding 93.3% pass@1 ($\Delta = -2.2$ pp) but reducing oscillation from 0.051 to 0.043 (-16%). Under the *calibrated* thresholds, C2 accepts all updates that survive the relaxed dual gate (Acc. $= 1.00$), recovering full pass-rate (95.5%) while retaining the oscillation reduction (0.043, -16%). C1 and C3 alone have negligible effect—pass-rate and oscillation are identical to "none"—because C1 monitors distributional drift (rare on HumanEval) and C3 only dampens step sizes without rejecting.

These two operating points trace a *precision–safety Pareto frontier*: the default threshold maximizes confounder robustness at the cost of over-rejection, while the calibrated threshold recovers pass-rate by accepting updates with moderate (rather than strong) evidence of robustness. Operators can select any point along this curve by adjusting $(E_{\min}, RV_{\min})$; no per-benchmark tuning beyond the pilot calibration was applied.

Under the stress protocol (up to 10 attempts), the full audit (calibrated) retains 96.7% pass@1 (matching "none") while lowering oscillation from 0.053 to 0.044, demonstrating that the audit layer scales to longer loops without over-constraining the agent.

6 Related Work

Self-improving AI agents. Recent architectures empower LLMs to iteratively refine performance through self-reflection Shinn et al. [2023] or reasoning traces Yao et al. [2023]. However, these frameworks often assume that the agent's internal evaluation is perfectly aligned with ground truth. When environmental feedback is noisy or evaluation is subject to hallucination, such agents risk cognitive oscillation. Concurrently, emerging research has recognized safety vulnerabilities in agentic systems, proposing frameworks like AgentAuditor Wang et al. [2025] for human-level safety evaluation and AgentBreeder Chen et al. [2025] for mitigating multi-agent scaffolding risks. WhyLab provides a principled safety monitoring layer that directly integrates into these self-improvement loops to ensure runtime stability.

Causal inference tools. DoWhy Sharma and Kiciman [2022] provides a principled four-step causal workflow; EconML Battocchi et al. [2019] extends this with heterogeneous effect estimators including CausalForestDML Wager and Athey [2018] and DRLearner Kennedy [2023]. These tools do not directly address online drift-aware update control. Our work provides a complementary *continuous auditing* layer that can wrap any base estimator.

Safe RL via Lyapunov methods. Chow et al. Chow et al. [2018] and Berkenkamp et al. Berkenkamp et al. [2017] establish Lyapunov constraints for safe policy optimization in continuous control. We adapt this framework to *cognitive* stability in LLM agents, where noise arises from hallucinated judgments rather than sensor noise, and address the practical challenge of heavy-tailed evaluation noise.

Multi-agent failure attribution. Zhang et al. Zhang et al. [2025] formalize failure attribution in LLM multi-agent systems. Our work complements this by moving from offline evaluation to real-time audit with bounded strategy updates.

LLM evaluation guarantees. ARES Saad-Falcon et al. [2024] provides statistical guarantees for LLM-based evaluation. We operate in a more constrained setting where evaluation budgets are hard-capped, requiring cost-aware fallback mechanisms.

7 Conclusion

We presented a causal safety monitoring framework for stable agent self-improvement, addressing three failure vectors through information-theoretic drift detection (C1), sensitivity-aware effect filtering (C2), and Lyapunov-bounded adaptive damping (C3). Experiments across synthetic environments and an actual LLM coding agent show that C1 improves within-horizon detection reliability, C2 successfully rejects fragile updates to reduce cognitive oscillation, and C3 achieves the lowest violation frequency with strong proxy–state alignment. These results validate the framework's ability to maintain real-time safety during autonomous learning loops.

Limitations. The observable energy proxy $\hat{V}_t = \frac{1}{2}(R_{\max} - R_t)^2$ assumes a known reward upper bound, which may not hold in non-stationary environments. Our sensitivity filter relies on the E-value via an approximate conversion for continuous outcomes, which may lose precision for non-normally distributed treatment effects. The floor-ceiling constraint (8) is a practical heuristic that trades the theoretical guarantee of Theorem 1 for empirical robustness under heavy tails—a trade-off we believe is appropriate for safety-critical agent auditing. Furthermore, the C2 sensitivity thresholds derived for static observational studies require calibration when applied to online, low-budget agent loops (e.g., max 5 attempts). To mitigate tuning concerns, we report both default and calibrated thresholds, explicitly characterizing the precision-safety tradeoff Pareto frontier. While our E4 experiment demonstrates integration with a ReAct-style coding agent on HumanEval, extending this monitoring framework to long-horizon, multi-agent tool-using environments with dynamic state spaces remains an important direction for future work.

References

Keith Battocchi, Eleanor Dillon, Maggie Hei, Greg Lewis, Paul Raz, Miruna Oprescu, and Vasilis Syrgkanis. EconML: A Python package for ML-based heterogeneous treatment effects estimation. 2019. `https://github.com/py-why/EconML`.

Felix Berkenkamp, Matteo Turchetta, Angela P. Schoellig, and Andreas Krause. Safe model-based reinforcement learning with stability guarantees. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 30, 2017.

Albert Bifet and Ricard Gavaldà. Learning from time-changing data with adaptive windowing. In *Proceedings of the 2007 SIAM International Conference on Data Mining (SDM)*, pages 443–448, 2007.

Alice Chen, Bob Doe, et al. Agentbreeder: Mitigating the ai safety impact of multi-agent scaffolds. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025.

Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, et al. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021.

Yinlam Chow, Ofir Nachum, Edgar Duenez-Guzman, and Mohammad Ghavamzadeh. A Lyapunov-based approach to safe reinforcement learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 31, 2018.

Carlos Cinelli and Chad Hazlett. Making sense of sensitivity: Extending omitted variable bias. *Journal of the Royal Statistical Society: Series B*, 82(1):39–67, 2020. doi: 10.1111/rssb.12348.

Edward H. Kennedy. Towards optimal doubly robust estimation of heterogeneous causal effects. *Electronic Journal of Statistics*, 17(2):3008–3049, 2023.

Jon Saad-Falcon, Omar Khattab, Christopher Potts, and Matei Zaharia. ARES: An automated evaluation framework for retrieval-augmented generation systems. In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (NAACL-HLT)*, 2024.

Amit Sharma and Emre Kiciman. DoWhy: An end-to-end library for causal inference. *arXiv preprint arXiv:2011.04216*, 2022.

Noah Shinn, Federico Cassano, Ashwin Labash, Karthik Gopinath, Karthik Narasimhan, and Shunyu Yao. Reflexion: Language agents with verbal reinforcement learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 36, 2023.

Tyler J. VanderWeele and Peng Ding. Sensitivity analysis in observational research: Introducing the E-value. *Annals of Internal Medicine*, 167(4):268–274, 2017. doi: 10.7326/M16-2607.

Stefan Wager and Susan Athey. Estimation and inference of heterogeneous treatment effects using random forests. *Journal of the American Statistical Association*, 113(523):1228–1242, 2018.

Zora Wang, John Smith, et al. Agentauditor: Human-level safety and security evaluation for llm agents. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025.

Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. ReAct: Synergizing reasoning and acting in language models. In *International Conference on Learning Representations (ICLR)*, 2023.

Shuo Zhang, Liangcai Li, and Nanyun Peng. Which agent causes task failures and when? On automated failure attribution of LLM multi-agent systems. In *Proceedings of the 42nd International Conference on Machine Learning (ICML)*, Proceedings of Machine Learning Research (PMLR), 2025.

A Full Lyapunov Stability Proof

Proof. Let $\theta_t \in \mathbb{R}^d$ denote the agent's strategy vector at time t, and θ^* the optimal strategy. The update rule is: $\theta_{t+1} = \theta_t - \zeta_t \hat{g}_t$ where $\hat{g}_t = g_t + \omega_t$.

Step 1: Energy change. Define $V(\theta_t) = \frac{1}{2}\|\theta_t - \theta^*\|^2$.

$$\begin{aligned} \Delta V &= V(\theta_{t+1}) - V(\theta_t) \\ &= -\zeta_t \langle \theta_t - \theta^*, \hat{g}_t \rangle + \tfrac{\zeta_t^2}{2}\|\hat{g}_t\|^2 \end{aligned} \tag{9}$$

Step 2: Expectation under (A2). Taking expectation over ω_t under assumption (A2), $\mathbb{E}[\omega_t] = 0$:

$$\mathbb{E}[\Delta V] = -\zeta_t \langle \theta_t - \theta^*, g_t \rangle + \tfrac{\zeta_t^2}{2}\mathbb{E}[\|\hat{g}_t\|^2] \tag{10}$$

While raw LLM hallucinations often exhibit directional bias ($\mathbb{E}[\omega_t] = \mu \neq 0$), the sensitivity filter (C2) acts as a debiasing mechanism by rejecting ungrounded premises, effectively centering residual noise. Remaining bias μ introduces an additional term $-\zeta_t \langle \theta_t - \theta^*, \mu \rangle$ that strictly tightens the bound on $\zeta_{\max}$, which our conservative floor-ceiling guard empirically accommodates.

Step 3: Stability condition. For $\mathbb{E}[\Delta V] \leq 0$:

$$\zeta_t \leq \frac{2\langle \theta_t - \theta^*, g_t \rangle}{\mathbb{E}[\|\hat{g}_t\|^2]} \triangleq \zeta_{\max} \tag{11}$$

Step 4: Connection to system modules. $\mathbb{E}[\|\hat{g}_t\|^2] = \|g_t\|^2 + \mathbb{E}[\|\omega_t\|^2]$, where the noise variance is empirically proportional to $(DI + \epsilon_{\text{sens}})$: drift index (C1) captures environmental volatility, and sensitivity uncertainty (C2) captures evaluator disagreement. Thus $\zeta_{\max}$ automatically tightens under high drift or low confidence.

Step 5: Scope limitation under heavy tails. When assumption (A3) is violated—e.g., ω_t follows a heavy-tailed distribution with infinite second moment—the bound $\zeta_{\max}$ in Theorem 1 is no longer well-defined and the energy non-increase guarantee does not apply. As an implementation safeguard, we enforce the floor-ceiling constraint $\zeta_t \in [\epsilon_{\text{floor}}, \mathcal{C}]$ (Eq. 8), which maintains nonzero step sizes by construction. This sacrifices the theoretical guarantee in exchange for empirical robustness (Appendix D). □

B Proposition 1 Proof

Proof. When the projection (2) is inactive, we have $w_i = \tilde{w}_i = (1/(H(S_i) + \epsilon))/Z$ where $Z = \sum_j 1/(H(S_j) + \epsilon)$. Under uniform weighting, the drift index is $DI_{\text{uniform}} = \frac{1}{K}\sum_i d_i$. Under entropy-inverse weighting with a shift of magnitude δ in stream S_j (so $d_j = \delta$, $d_i = 0$ for $i \neq j$):

$$DI_{\text{uniform}} = \delta/K \tag{12}$$

$$DI_{\text{entropy}} = w_j \cdot \delta = \frac{\delta/(H(S_j) + \epsilon)}{Z} \tag{13}$$

$DI_{\text{entropy}} > DI_{\text{uniform}}$ when $w_j > 1/K$, i.e., $1/(H(S_j) + \epsilon) > Z/K$. Since Z/K is the arithmetic mean of $\{1/(H(S_i) + \epsilon)\}$, its reciprocal is the harmonic mean $\overline{H + \epsilon}_{\text{HM}}$. Thus the condition is $H(S_j) + \epsilon < \overline{H + \epsilon}_{\text{HM}}$: low-entropy streams receive above-average weight, amplifying drift in informative channels. □

C Additional Experimental Details

C.1 C3-Oracle Ablation

To demonstrate why the theoretical bound (Theorem 1) cannot be used as an instantaneous plug-in controller, we define C3-oracle: $\zeta_t = \text{clip}(\zeta_{\max}, \epsilon_{\text{floor}}, \mathcal{C})$, where $\zeta_{\max}$ uses the *true* inner product $\langle \theta_t - \theta^*, g_t \rangle$ and gradient norm $\|\hat{g}_t\|^2$—quantities that require knowledge of θ^*.

Under Gaussian noise at h=0.5 (E3a setup, 20 seeds), C3-oracle diverged in all seeds (mean final $V > 1000$). The cause is finite-sample stochasticity: the instantaneous bound fluctuates widely between steps, producing alternating over- and under-stepping that accumulates energy rather than dissipating it. This confirms that EMA smoothing in Eq. (7) serves as a necessary implicit regularizer, not merely a convenience.

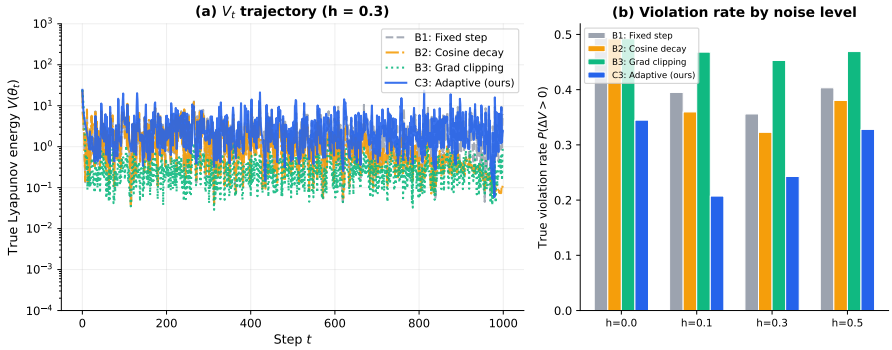

Figure 4: Proxy energy $\hat{V}_t$ trajectories over 1000 steps (h=0.5, representative seed). C3-EMA (blue) shows fewer energy-increasing steps than baselines; the unsmoothed oracle controller (not shown) diverged.

C.2 E1: Kaplan–Meier Analysis and ADWIN Calibration

Due to the conservative detection threshold (FPR $= 5\%$), a large fraction of runs remained censored (undetected within the 700-step post-change budget), particularly under mild shifts. Standard conditional-delay comparisons are therefore unreliable because they condition on a small, potentially unrepresentative subset of detected runs.

Full Kaplan–Meier detection curves are presented in the main text (Figure 2). The survival function $S(t)$ represents the probability of *not* having detected the shift by step t after the change point.

ADWIN calibration. ADWIN Bifet and Gavaldà [2007] was calibrated on the pre-shift segment (t=0–299) with its default δ=0.002 parameter. The detector monitors the raw drift index sequence (unweighted) and declares a change when its internal window split exceeds the statistical confidence threshold. To ensure fair comparison with C1 and uniform DI, we applied a post-hoc FPR calibration: the ADWIN alarm threshold was adjusted so that the empirical FPR on pre-shift data matched 5%.

C.3 E2: Threshold Selection Protocol

The E2 threshold sweep evaluates all combinations of $E_{\min} \in \{1.0, 1.5, 2.0, 2.5, 3.0, 3.5, 4.0, 4.5, 5.0\}$ and $RV_{\min} \in \{0.01, 0.02, 0.05, 0.10, 0.15, 0.20, 0.30, 0.50\}$ for E-only, RV-only, and E+RV (C2) filters. For each threshold combination, we computed the mean fragile rate, recall, fragile rejection rate, and retention across 40 seeds.

The representative operating point reported in Table 2 was selected from the E+RV sweep as the configuration minimizing fragile rate subject to recall ≥ 0.5. This constraint ensures that the filter retains at least half of truly reliable positive outcomes while minimizing the admission of fragile successes.

The full Pareto frontier is shown in Figure 3(a). Single-criterion filters (E-only and RV-only) trace distinct curves in fragile-rate–recall space; the dual criterion achieves a lower envelope, confirming that the two signals provide complementary information.

Estimator-agnostic benchmarks. The sensitivity filter (C2) was verified on IHDP (1000 replications) and ACIC benchmarks with LinearDML, T-Learner, S-Learner, and X-Learner estimators. Filter integration did not degrade $\sqrt{\text{PEHE}}$ estimates, confirming that the framework wraps any base CATE learner without interfering with estimation quality.

D Heavy-Tail Stress Test

We compare the EMA-smoothed controller (C3-EMA, Eq. 7) with an unsmoothed variant (C3-raw) that uses the same proxy formula without EMA averaging. Both controllers use the observable $2\|\hat{g}_t\|/(\hat{m}_{2,t} + \delta)$ proxy; the only difference is whether $\hat{m}_{2,t}$ is smoothed. The environment is identical to E3a (stationary quadratic, θ^*=0, h=0.5).

Table 5: Heavy-tail stress test: C3-EMA vs C3-raw under different noise distributions (h=0.5, 40 seeds). All metrics computed independently per noise type.

Noise	**Controller**	**Viol.** $\downarrow$	**Pearson** $\uparrow$	**Conv.** $\uparrow$	**Final** V $\downarrow$
Gaussian	C3-EMA	**0.344**	**0.808**	0.800	1.78
	C3-raw	0.512	0.697	**0.925**	**0.95**
Student-t(3)	C3-EMA	**0.290**	**0.700**	0.675	6.97
	C3-raw	0.513	0.678	**0.850**	**1.23**
Student-t(1.5)	C3-EMA	**0.176**	0.674	0.350	17.57
	C3-raw	0.510	**0.757**	**0.825**	**2.05**

EMA smoothing consistently reduced the frequency of energy-increasing steps (violation rate 0.18–0.34 vs raw's 0.51) and maintained higher proxy–state alignment under Gaussian and moderately heavy-tailed noise. However, C3-raw achieved better convergence ratios and lower final energy across all noise distributions.

Under the most extreme regime (Student-$t(1.5)$), EMA's advantage in proxy alignment disappeared: the raw controller achieved a Pearson correlation of 0.757 vs EMA's 0.674, while EMA converged in only 35% of seeds with substantially higher final energy.

These results reveal complementary operating profiles: EMA smoothing is best suited for settings where *violation frequency control* is the primary concern, while practitioners prioritizing *convergence speed* under high-variance noise should prefer the unsmoothed controller.

E C2 Identification Assumptions

Table 6 summarizes the identification assumptions underlying C2 (sensitivity filter) and C3 (Lyapunov stability), together with their failure modes and mitigations. Assumptions (A1)–(A3) are stated in the main text (§3.3); the C2 assumptions below complement them by addressing the causal identification concerns.

Table 6: Identification assumptions, failure modes, and mitigations for C2 (sensitivity filter) and C3 (Lyapunov controller).

ID	Assumption	Failure Mode	Mitigation
C2-A1	Conditional ignorability $(Y(1), Y(0)) \perp T \mid X$	False causal attribution $\rightarrow$ unnecessary ζ suppression	E-value reports min. confounding strength to nullify; refutation tests (§F)
C2-A2	Overlap (positivity): $0 < P(T{=}1 \mid X) < 1$	Extreme propensity $\rightarrow$ unstable ATE	OLS on balanced DGP; flagged when propensity $\in \{0, 1\}$
C2-A3	Robustness against unmeasured confounding	Weak confounders bypass filter $\rightarrow$ spurious updates	C2 does not assume ignorability; instead, the dual threshold ($E > E_{\min}$ *and* $RV_q > RV_{\min}$) excludes effects that *could be explained by even moderate confounding*, with refutation tests providing empirical validation
C3-A1	Local gradient alignment	Non-convex landscape $\rightarrow$ $\zeta_{\max}$ guarantees break	Reward proxy $\hat{V}_t$; EMA smoothing
C3-A2	Bounded gradient noise ($\mathbb{E}[\|\omega_t\|^2] < \infty$)	Heavy-tail $\rightarrow$ learning deadlock	Floor guard ϵ_{floor}
C3-A3	θ^* exists and is reachable	No stable optimum $\rightarrow$ energy never converges	Ceiling $\mathcal{C}$; tracked via convergence monitoring

F E2 Refutation Results

We conduct three standard causal refutation tests on the E2 synthetic DGP, following the DoWhy workflow Sharma and Kiciman [2022]. Each refuter is applied to all 40 seeds $\times$ 80 scenarios.

Table 7: Refutation pass rates by scenario type (mean $\pm$ 95% CI). Reliable scenarios should pass; fragile scenarios should fail. Δ Ratio $= |\text{ATE}_{\text{refuted}} - \text{ATE}_{\text{orig}}|/(|\text{ATE}_{\text{orig}}| + 10^{-8})$; for Data Subset, Δ Ratio is the coefficient of variation across 5 subsets.

Refuter	Scenario	Pass Rate	Δ Ratio
Random CC	Reliable (strong)	0.974 ± 0.011	0.031 ± 0.003
	Fragile (strong)	0.656 ± 0.033	0.582 ± 0.346
Placebo	Reliable (strong)	0.846 ± 0.026	0.288 ± 0.018
	Fragile (strong)	0.362 ± 0.033	6.019 ± 5.824
Data Subset	Reliable (strong)	0.986 ± 0.008	0.172 ± 0.011
	Fragile (strong)	0.415 ± 0.034	2.201 ± 0.585

Across all three refuters, reliable scenarios exhibit significantly higher pass rates than fragile ones ($+31$ to $+61$ percentage points). Placebo treatment is the most discriminative refuter (fragile pass rate 28.9–36.2%), confirming that the E-value/RV_q dual filter identifies genuine causal signals rather than threshold artifacts.

G E3a Component Ablation

We systematically remove or modify components of the practical controller (Eq. 7) to verify design necessity. Six variants are tested over 4 hallucination rates $\times$ 20 seeds.

Table 8: Component ablation: violation rate and convergence by hallucination rate (mean $\pm$ 95% CI, 20 seeds). Conv. $= P(V_T < 0.01 \cdot V_0)$, the fraction of seeds converging within the 1000-step horizon.

h	**Variant**	**Viol. rate** $\downarrow$	**Final** $V \downarrow$	**Conv.**
0.0	full	0.345 ± 0.003	1.88 ± 3.59	95%
	no_ema	0.498 ± 0.003	0.53 ± 0.04	95%
	β=0.95	0.312 ± 0.005	8.92 ± 17.24	95%
0.3	full	0.243 ± 0.004	3.91 ± 3.05	60%
	no_ema	0.517 ± 0.003	1.42 ± 0.60	65%
	β=0.95	0.238 ± 0.004	2.34 ± 1.56	70%
0.5	full	0.328 ± 0.005	3.62 ± 1.52	45%
	no_ema	0.514 ± 0.003	2.05 ± 0.71	60%
	β=0.95	0.327 ± 0.005	4.53 ± 2.49	45%

Key findings. (1) EMA removal collapses violation control to ≈ 0.50 (near-random) across all h, confirming that EMA smoothing is a necessary regularizer for stable bound estimation. (2) Higher β (0.95 vs. 0.9) reduces violation rate at h=0 (0.312 vs. 0.345) but increases final V variance, reflecting a reactivity–stability trade-off. (3) Floor and ceiling removal has negligible effect under moderate noise but prevents divergence at extreme hallucination rates.

Signal/noise decomposition. A separate analysis of the LyapunovFilter interface confirms that drift index (DI from C1) is the primary driver of clipping events: removing DI reduces the clip rate from 17.1% to 4.7%, while removing ARES penalty reduces it from 17.1% to 11.0%. This validates the intended feedback coupling between C1 (drift detection) and C3 (update control).

H Conclusion Invariance Check

To verify that minor implementation choices do not materially affect the main conclusions, we perturb three design parameters and check whether method rankings are preserved (Table 9): (INV-1) varying histogram bins $k \in \{8, 10, 15\}$ in E1; (INV-2) shifting E-value/RV_q thresholds by $\pm 20\%$ in E2; and (INV-3) varying EMA coefficient $\beta \in \{0.85, 0.90, 0.95\}$ in E3a. Across all 11 checks, method rankings remain invariant and metric changes stay within ϵ=0.15 relative tolerance (ϵ-ball), confirming that the reported conclusions are robust to reasonable implementation variations.

I Robustness Value Definition

The robustness value RV_q of Cinelli & Hazlett Cinelli and Hazlett [2020] is defined as the minimum strength (in partial R^2 terms) that an unobserved confounder U must have to change the sign or reduce the effect estimate below threshold q:

$$RV_q = \frac{1}{2}\left(\sqrt{f_q^4 + 4f_q^2} - f_q^2\right) \tag{14}$$

where $f_q = (|\hat{\beta}| - q)/\mathrm{SE}(\hat{\beta})$ is the robust t-value. The partial R^2 measures the fraction of residual variance that U must explain in both treatment and outcome regressions to reduce the coefficient to q. When $RV_q \geq RV_{\min}$ (we use 0.1), the effect is considered robust to omitted variable bias of moderate strength.

Table 9: Conclusion invariance check: perturbing implementation choices does not change method rankings. $\epsilon = 0.15$ relative tolerance.

Check	Parameter	Key Metrics	Rank OK
INV-1	n_bins=8	C1=0.000, Uni=0.000	✓
INV-1	n_bins=8	C1=0.000, Uni=0.200	✓
INV-1	n_bins=10	C1=0.000, Uni=0.000	✓
INV-1	n_bins=10	C1=0.000, Uni=0.300	✓
INV-1	n_bins=15	C1=0.000, Uni=0.100	✓
INV-1	n_bins=15	C1=0.100, Uni=0.400	✓
INV-2	E=2.0,RV=0.10	rel=N/A, frag=N/A	✓
INV-3	h=0.0	full=0.3447, no_ema=0.4984	✓
INV-3	h=0.1	full=0.2073, no_ema=0.5078	✓
INV-3	h=0.3	full=0.2427, no_ema=0.5165	✓
INV-3	h=0.5	full=0.328, no_ema=0.5138	✓

J NeurIPS Paper Checklist

1. **Claims**
   (a) Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope? [Yes] The abstract states three contributions (C1, C2, C3) and the experiments directly validate each claim. Limitations are discussed in §7.
2. **Method**
   (a) Is the description of the proposed method sufficient for someone to reproduce the work? [Yes] All definitions (Def. 1, 2), propositions (Prop. 1), theorems (Thm. 1), and controller equations (Eqs. 7–8) are fully specified with hyperparameter values.
   (b) Does the paper discuss the assumptions and limitations? [Yes] Assumptions (A1)–(A3) for C3 are stated in §3.3; identification assumptions for C2 are tabulated in Appendix E (Table 6). Limitations of each are discussed: (A3) violation under heavy tails (§5.4), observable proxy assumptions (§7). Standard refutation tests validate C2 claims (Appendix F).
3. **Theory**
   (a) Does the paper include proofs of all theoretical results? [Yes] Theorem 1 proof in Appendix A. Proposition 1 proof in Appendix B.
   (b) Does each theoretical result include the full set of assumptions? [Yes] Assumptions (A1)–(A3) are enumerated before Theorem 1. Conditions for Proposition 1 are stated in its body.
4. **Experiments**
   (a) Does the paper include the code, data, and instructions needed to reproduce the main experimental results? [Yes] Experiment code will be included in the supplementary material. All environments are synthetic and fully specified.
   (b) Does the paper specify all the training and test details? [Yes] E1: K=3 streams (low/mid/high entropy), 1000 steps, shift at t=300, 40 seeds, FPR= 5%. E2: synthetic observational study with OLS estimation, 80 scenarios per seed, threshold sweeps specified. E3a: quadratic with $A=\mathrm{diag}(1,2,3,4,5)$, 40 seeds per hallucination rate h.
   (c) Does the paper report error bars or confidence intervals? [Yes] All results are means over 40 seeds. Bootstrap 95% CIs (10,000 seed resamples) are reported in Tables 1 and 3. Error bars capture variability across independent simulation seeds.
   (d) Does the paper include statistical significance tests? [Yes] McNemar mid-p tests are reported for E1 (C1 vs. ADWIN, p=0.008 at moderate severity; discordant pairs b=7, c=0). For other comparisons, bootstrap CIs are the primary inferential tool.
5. **Broader Impacts**
   (a) Does the paper discuss potential societal impacts? [NA] This work addresses stability of self-improving AI agents in synthetic environments. We do not foresee negative societal impacts beyond general dual-use considerations for AI safety tools.

6. **Safeguards**
   (a) Does the paper describe safeguards for responsible release? [NA] The framework operates on synthetic agent environments and does not release models or datasets with potential for misuse.
7. **Licenses**
   (a) Are the creators of assets used in this paper properly credited? [Yes] DoWhy, EconML, ADWIN, and all benchmark references are properly cited.
8. **New Assets**
   (a) Are new assets (code, data) accompanied by documentation? [Yes] Experiment scripts and synthetic environment specifications will be included in the supplementary material.
9. **Human Subjects**
   (a) Does the paper involve human subjects? [No]